

# Final Report

## 1. Overview

This project uses PyTorch and the ResNet18 architecture to classify images. The process begins with loading and concatenating `.npy` files containing data and labels, followed by creating a custom dataset, applying transformations/augmentations, and finally training and validating a model. Below is a detailed explanation of each step in the workflow.

## 2. Data Import and Concatenation

- **Data Files and Label Files:** The script specifies multiple `.npy` files for the images (`data0.npy`, `data1.npy`, `data2.npy`) and corresponding label files (`lab0.npy`, `lab1.npy`, `lab2.npy`).
- **Loading and Merging:** Each file is loaded with NumPy's `np.load`. The resulting arrays are concatenated along the first axis. This results in two large NumPy arrays: one containing all the images ( `data` ) and the other containing all the labels ( `labels` ).
- **Shape Verification:** After concatenation, the script prints out the shapes of the `data` and `labels` arrays to confirm that everything loaded correctly.

## 3. Custom Dataset Class

- **Purpose:** A `CustomDataset` class, extending PyTorch's `Dataset`, is defined to handle the data in a convenient way for `DataLoader`.
- **Key Steps:**
  1. The dataset stores `data` and `labels`.
  2. The `__getitem__` method:

- Reshapes the sample from (40, 168) to maintain its single-channel (grayscale) format.
  - Converts it to float32.
  - Optionally applies a transformation pipeline if provided.
3. The `__len__` method returns the total number of samples in the dataset.

## 4. Data Transformations (Augmentations)

- **Training Transform:**

- Converts the NumPy array to a PIL image.
- Resizes it to 224×224 (matching ResNet18 input size).
- Applies random horizontal flips, random rotations, and random affine transformations to augment the dataset.
- Converts the image back to a PyTorch tensor and normalizes pixel values (mean=0.5, std=0.5).

- **Validation Transform:**

- Similar to the training transform, but typically without random augmentations (so the evaluation data remains consistent).

These transformations help improve generalization by exposing the model to a variety of augmented samples during training.

## 5. Splitting the Dataset

- **Train, Validation, and Test Sets:**

- The first 20,000 samples are designated for training, with full augmentations.

- The next 4,000 samples (from index 20,000 to 24,000) are for validation, with a simpler transform (no random augmentations).
- All remaining samples (index 24,000 onwards) form the test set, often with either no transform or only basic normalization.

## 6. DataLoaders

- **Creation:** Three DataLoaders ( `train_loader` , `val_loader` , and `test_loader` ) wrap the custom datasets.
- **Parameters:**
  - Batch size is 32 for all sets.
  - Training data is shuffled at every epoch ( `shuffle=True` ), while validation and test sets are not shuffled ( `shuffle=False` ).

## 7. Model Definition (ResNet18)

- **Pretrained Model:** A pre-trained ResNet18 from `torchvision.models` is loaded with `pretrained=True` , which initializes weights from ImageNet.
- **Input Channel Modification:** The original ResNet expects 3 input channels (RGB). Here, the first convolution layer is replaced to accept a single grayscale channel.
- **Output Layer Modification:** The final fully connected layer ( `fc` ) is replaced with a new layer whose output size matches the number of classes in the dataset. This is determined by the unique labels.

## 8. Loss Function, Optimizer, and Scheduler

- **Loss Function:** `nn.CrossEntropyLoss` is used for multi-class classification.

- **Optimizer:**An Adam optimizer is chosen, starting with a learning rate of 0.001.
- **Scheduler:**A StepLR scheduler reduces the learning rate by a factor ( `gamma=0.1` ) every 5 epochs (set by `step_size=5` ).

## 9. Training Loop

- **Function:**A function `train_model` encapsulates the training/validation procedure.
- **Training Phase:**
  1. `model.train()` is called to set the model to training mode.
  2. For each batch, images and labels are moved to the device (GPU/CPU).
  3. The optimizer gradients are zeroed.
  4. A forward pass computes predictions, and the cross-entropy loss is calculated.
  5. A backward pass ( `loss.backward()` ) updates the model's weights via `optimizer.step()` .
  6. Accumulates training loss and the number of correct predictions.
- **Validation Phase:**
  1. `model.eval()` is called, and gradients are not computed ( `with torch.no_grad()` ).
  2. Validation loss and accuracy are computed over the validation set.
- **Scheduler Step:**The learning rate is updated at the end of each epoch via the scheduler.
- **Saving the Best Model:**The best validation accuracy observed across epochs is stored, and the model's weights are saved to `best_model.pth` if an improvement occurs.

## 10. Model Training and Future Testing

- **Model Training:**The script trains for 10 epochs. During each epoch, it prints training/validation losses and accuracies.
- **Model Testing :**A test model function could then be applied to the test set to measure final performance. This step is commented out in the code, but can be implemented similarly to validation.